

# Advanced NLP - The Road to LLMs

Prof. Dr. Richard Sieg  
TH Köln IWS - WS 25/26

# Goal of this Course

Insights into the development of the first LLMs and  
their basic architecture & design.

How to use and steer LLMs programmatically for  
research & industry.

Sensitivity for shortcomings of LLMs: Bias,  
hallucinations...



With great power  
comes great responsibility.

# Agenda

01. Organization

05. Appendix

02. NLP Recap

03. Text Simplification

04. Set - Up & Tutorial

01.

# Organization



## About Me

- ▶ First term as full-time Professor for NLP & LLMs
- ▶ Since 2023 external lecturer for NLP @ IWS
- ▶ Mathematics Background
- ▶ NLP & ML Engineer for eight years
- ▶ Many years of university level teaching
- ▶ [richardsieg.github.io](https://richardsieg.github.io)

# Shall we meet via Zoom?



In January I would like to have at least two extended sessions  
(One being the presentations Jan 15th)

# Schedule

Lecture	Date	Topic
1	03.12.2025	Introduction & NLP Recap
2	04.12.2025	RNNs and LSTMs
3	10.12.2025	Attentions & Transformers
4	11.12.2025	Encoder & Decoder Models
5	17.12.2025	Hackathon / Check-In
6	18.12.2025	LLM Architecture
7	07.01.2026	LLM Engineering
8	08.01.2026	Hackathon / Check-In
9	14.01.2026	LLM Shortcomings
10	15.01.2026	Final Presentations

# Requirements & Grading

The exam and grade consists of

- Code implementation of all discussed milestones (30%, per team)
  - Code needs to be functioning and explained via comments
- A short (max 2 pages) report (pdf format) on the implementation, results and the work done by the team members (20%, per team)
- A 10 minute oral presentation per team member on their individual work on the 15th of January 2026 (50%, individual)

The final code and report has to be ready in your team's GitHub Repo by the 16th of January 2026, 6pm.

# Presentation

- Oral Presentation of the code and evaluation results per individual team member
- No slides, just code
- Approx 10 minutes per presentation
- Split up your team in a meaningful way
- Small Q&A after each presentation

Find teams and fill out the form



# Teams

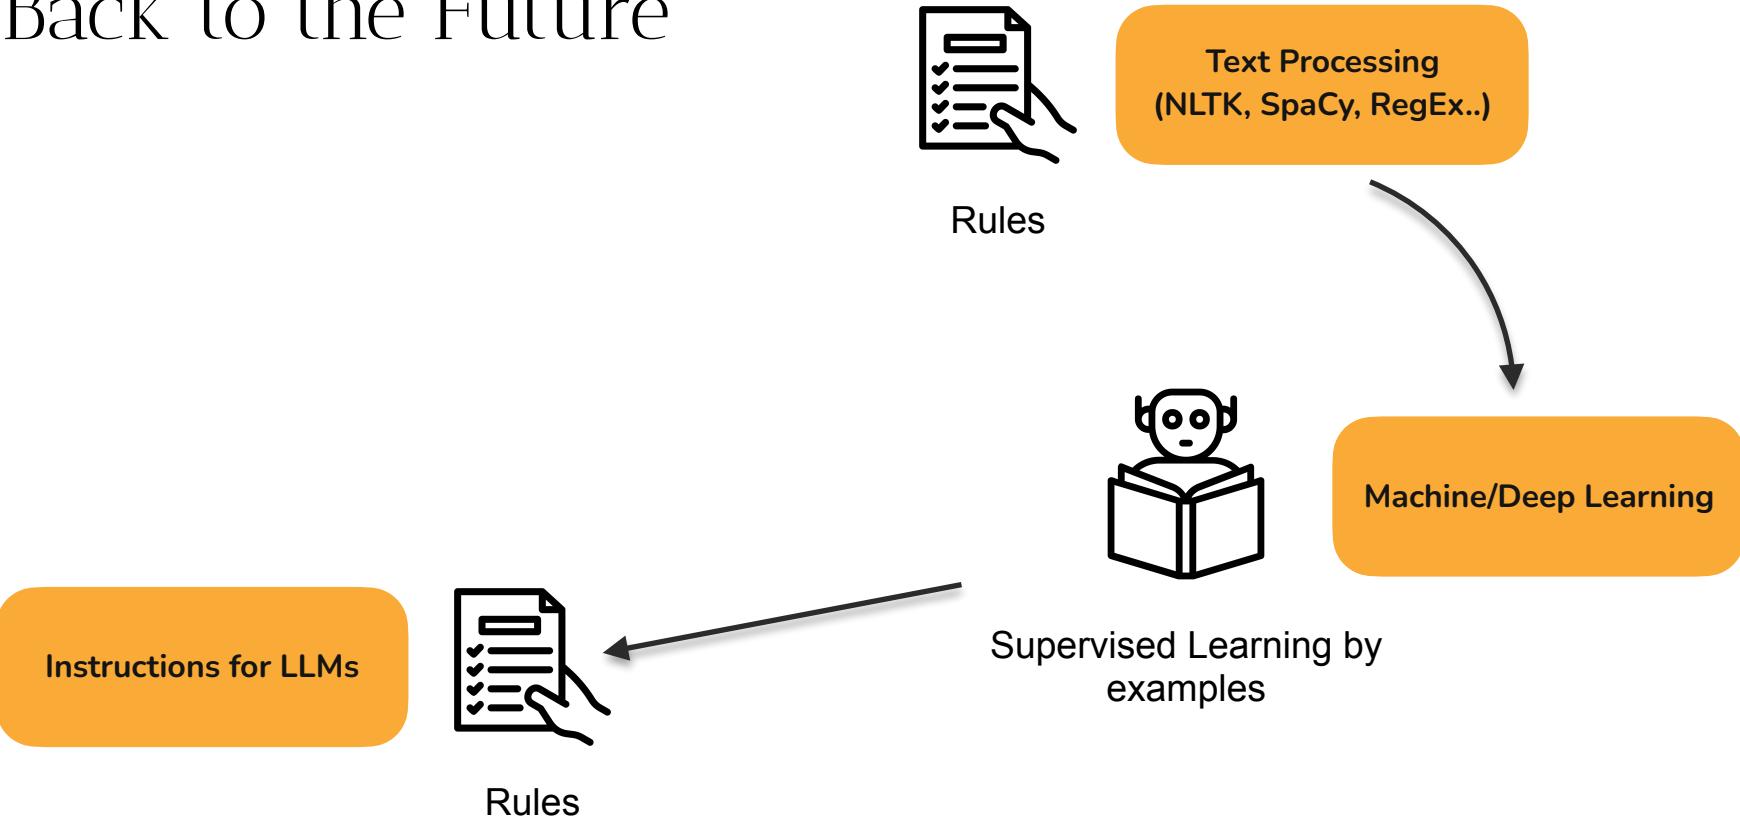
- Team 1
- Team 2
- Team 3
- Team 4
- ▶ Team 5
- ▶ Team 6
- ▶ Team 7
- ▶ Team 8

## 02. NLP Recap

# NLP so far

- Text Processing (spaCy, RegEx)
- Relations and semantics
- Frequency Based Embeddings (BoW, tf-idf)
- Prediction Based Word Embeddings (word2vec)
- Information extraction
- Text Classification
- Language Modeling

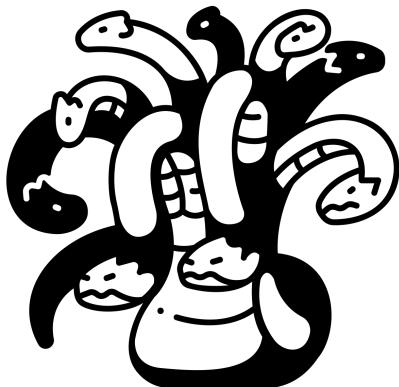
# Back to the Future



# The many-headed AI Hydra

One particular system or model

Large Language Models



“AI” means...

... and gets confused with  
Machine Learning

The research field AI

Artificial General  
Intelligence

You shall know a word by  
the company it keeps.

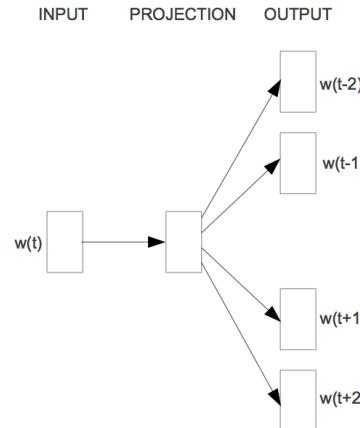
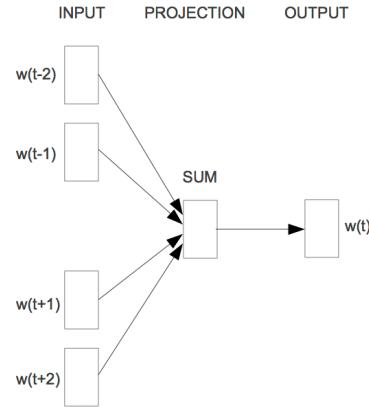
J.R. Firth 1957

# Frequency Based Embeddings

- Given: A number of documents
- **Bag of Words:** Count how many times a word appears in a document
  - Vector representation of one document
  - Vector dimension: size of vocabulary (unique words)
- **Term Frequency -~~Inverse Document Frequency~~<sub>|D|</sub>**  
$$tf(t, d) = \frac{f_{t,d}}{\sum_{t' \in d} f_{t',d}}$$
$$idf(t, D) = \log \frac{|D|}{|\{d \in D : t \in d\}|}$$
$$tfidf(t, d, D) = tf(t, d) \cdot idf(t, D)$$

# Predictive Word Embeddings

- Learn a vector representation of a word by predicting...
  - the word itself based on its context (Continuous Bag of Words)
  - the context words based on a target word (Skip-Gram)
- Word2vec, fastText, GloVe
- Issue: generates fixed embeddings, neural network itself cannot be used for further task specific fine-tuning



# Discriminative and Generating NLP Tasks

## Discriminative Tasks

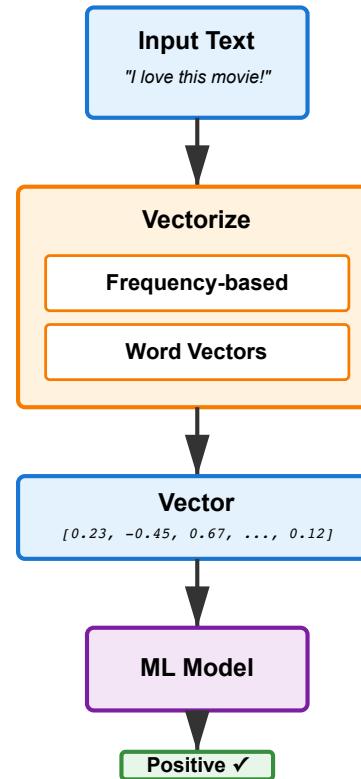
- Text Classification (e.g. sentiment)
- Named Entity Recognition
- Part of Speech Tagging
- Semantic Role Labeling

## Generative Tasks

- Chatbots
- Text Completion
- Machine Translation
- Question Answering
- Summarization
- **Simplification**

# Text Classification (so far)

- Given an input text vectorize it by
  - using frequency based embeddings
  - or combining the individual word vectors (e.g. average)
- Take this vector as an input for a classical ML model and train
  - Logistic regression, SVM, Naive Bayes, NN ...



# Language Modeling

- A *language model* is a statistical model that assigns a probability to a sequence of words (more precisely *tokens*)  
 $p(w_1, \dots, w_m)$
- Related task: Compute the probability of the next token given a sequence of previous tokens

$$p(w_t | w_{t-1}, \dots, w_{t-n+1})$$

- *Chain Rule*

$$p(w_1, \dots, w_m) = \prod_{i=1}^m p(w_i | w_1, \dots, w_{i-1})$$

- $p(\text{be water my friend}) = p(\text{be}) * p(\text{water} | \text{be}) * p(\text{my} | \text{be water}) * p(\text{friend} | \text{be water my})$

# Markov Assumption

- We can approximate the conditional probability by only look at the past n tokens (n-gram model)

- $\circ p(\text{friend} \mid \text{be water my}) \approx p(\text{friend} \mid \text{my}) \text{ (bigram)}$

- Bigram model: Predict the next word by only looking at the previous word

$$p(w_i \mid w_1, \dots, w_{i-1}) \approx p(w_i \mid w_{i-1})$$

- Calculation is based on counting frequencies

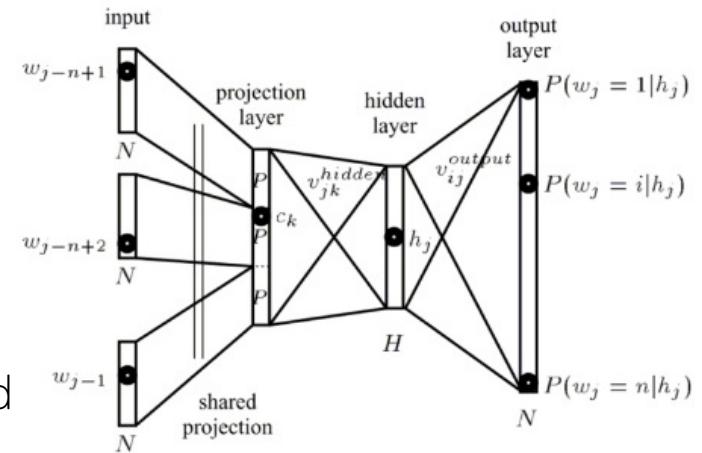
$$p(w_t \mid w_{t-1}, \dots, w_{t-n-1}) = \frac{\text{count}(w_{t-n+1}, \dots, w_{t-1}, w_t)}{\text{count}(w_{t-n+1}, \dots, w_{t-1})}$$

# Language Modeling with Neural Networks

- The conditional probability can be expressed as a softmax layer inside a neural network

$$p(w_t | w_{t-1}, \dots, w_{t-n-1}) = \frac{\exp(h^T v'_{w_t})}{\sum_{w_i} \exp(h^T v'_{w_i})}$$

- $h$ : context vector (output of the hidden layer on  $w_{t-1}, \dots, w_{t-n-1}$ )
- $v'_{w_t}$ : the embedding to be learned for the token/word  $w_t$
- Objective function: Maximize probability for the correct word



# History of Language Models

~1950s-1990s



## Statistical LMs

N-grams and Markov models

~Early 2000s



## Feedforward Neural LMs

Neural networks,  
Learned word  
embeddings jointly.

2013



## Word2Vec

Efficient dense word  
embeddings enable  
new applications

~2014-2017



## RNNs/LSTMs

Recurrent networks  
improve sequence  
handling

2017



## Transformer

Attention  
mechanism  
revolutionizes  
language modeling

2018-Present



## Large Language Models

BERT, GPT, T5, etc.

NLP so far

This lecture

# 03. Text Simplification

# Text Simplification

- The task of simplifying text while preserving meaning.
- Can be general or context specific, e.g. “fourth graders” or “Leichte Sprache” or “A1 reading level”.
- Before LLMs
  - Fine-Tuned Sequence to Sequence models
  - Word tagging and replacement
- Out-of-the-box LLMs already match or are superior to these approaches.
- ➔ Major parts of research is focused on utilizing LLMs nowadays.

# Measuring Simplicity

- FKGL: US Navy based score that measures relation of syllables, words, and sentences
- SARI: Needs source and reference sentence and measures the suitability of words that were added, deleted, and kept
- BLEU: Score from Machine Translation; measures syllable overlap with reference sentence
- BERTScore: Score from Machine Translation; compares the BERT sentence embeddings of simplification and reference
- SLE: RoBERTa model fine-tuned to score simplicity

## Source

Owls are the order Strigiformes, comprising 200 bird of prey species.

## Reference Simplification

An owl is a bird. There are about 200 kinds of owls.

## Predicted Simplification

Owls are a group of birds with about 200 species that hunt for food.

# The challenge

04. Set - Up

# uv

- New Package Manager that has replaced all others (pip, poetry etc) and is very fast
- Management of packages in a **pyproject.toml**
- Python Installation: **uv python install 3.12**
- First Installation and venv creation: **uv sync**
- Run scripts: **uv run ...**
- Add package: **uv add ...**

# marimo

- New alternative to Jupyter Notebooks
- Lives in Python files (-> better Diffs & Imports)
- Reactive: Dependent cells are automatically executed
- Interactive: Many possible Input Options
- Coding Assistant via GitHub Copilot & LLM APIs (OpenAI etc)
- Editor: In Browser or VSCode/Cursor Extension
- Start editor: **marimo edit**
- Run as Web App: **marimo run notebook.py**

# GitHub

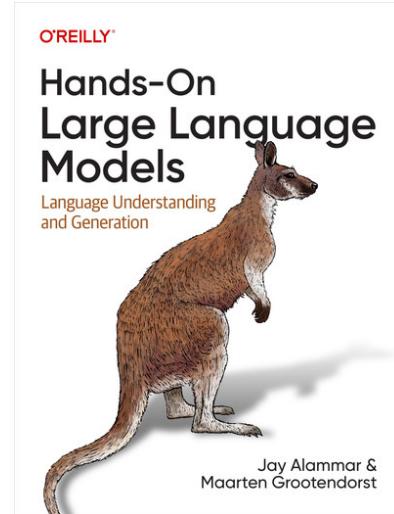
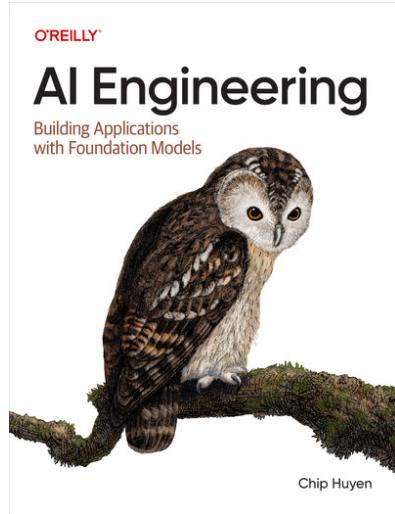
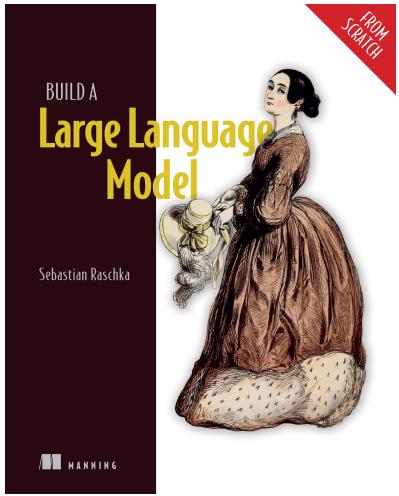
- One main GitHub repository for all slides and notebooks
- Please do a fork, invite your team members and me, and rename it so it contains your team number
- When you are all set, run **uv sync** and **uv run marimo edit**
- Remember to register for GitHub Education (Copilot Pro for free)

## 05. Appendix



**rhAIInland Community - Join our Discord!**

# Some Literature



# Folks to follow

- Godfathers of AI: Geoffrey Hinton, Yoshua Bengio, Yann LeCun
- Andrew Ng - Coursera founder, influential AI figure, has newsletter (The Batch)
- Andrej Karpathy - LLM guru, great YT videos, former OpenAI, now Tesla
- Chip Huyen
- Jay Alammar
- Sebastian Raschka
- Vincent Warmerdam - YouTuber, great talks, former spaCy, now marimo
- Ines Montani - Founder of spaCy
- *For more newsletters and blogs see richardsieg.github.io*